

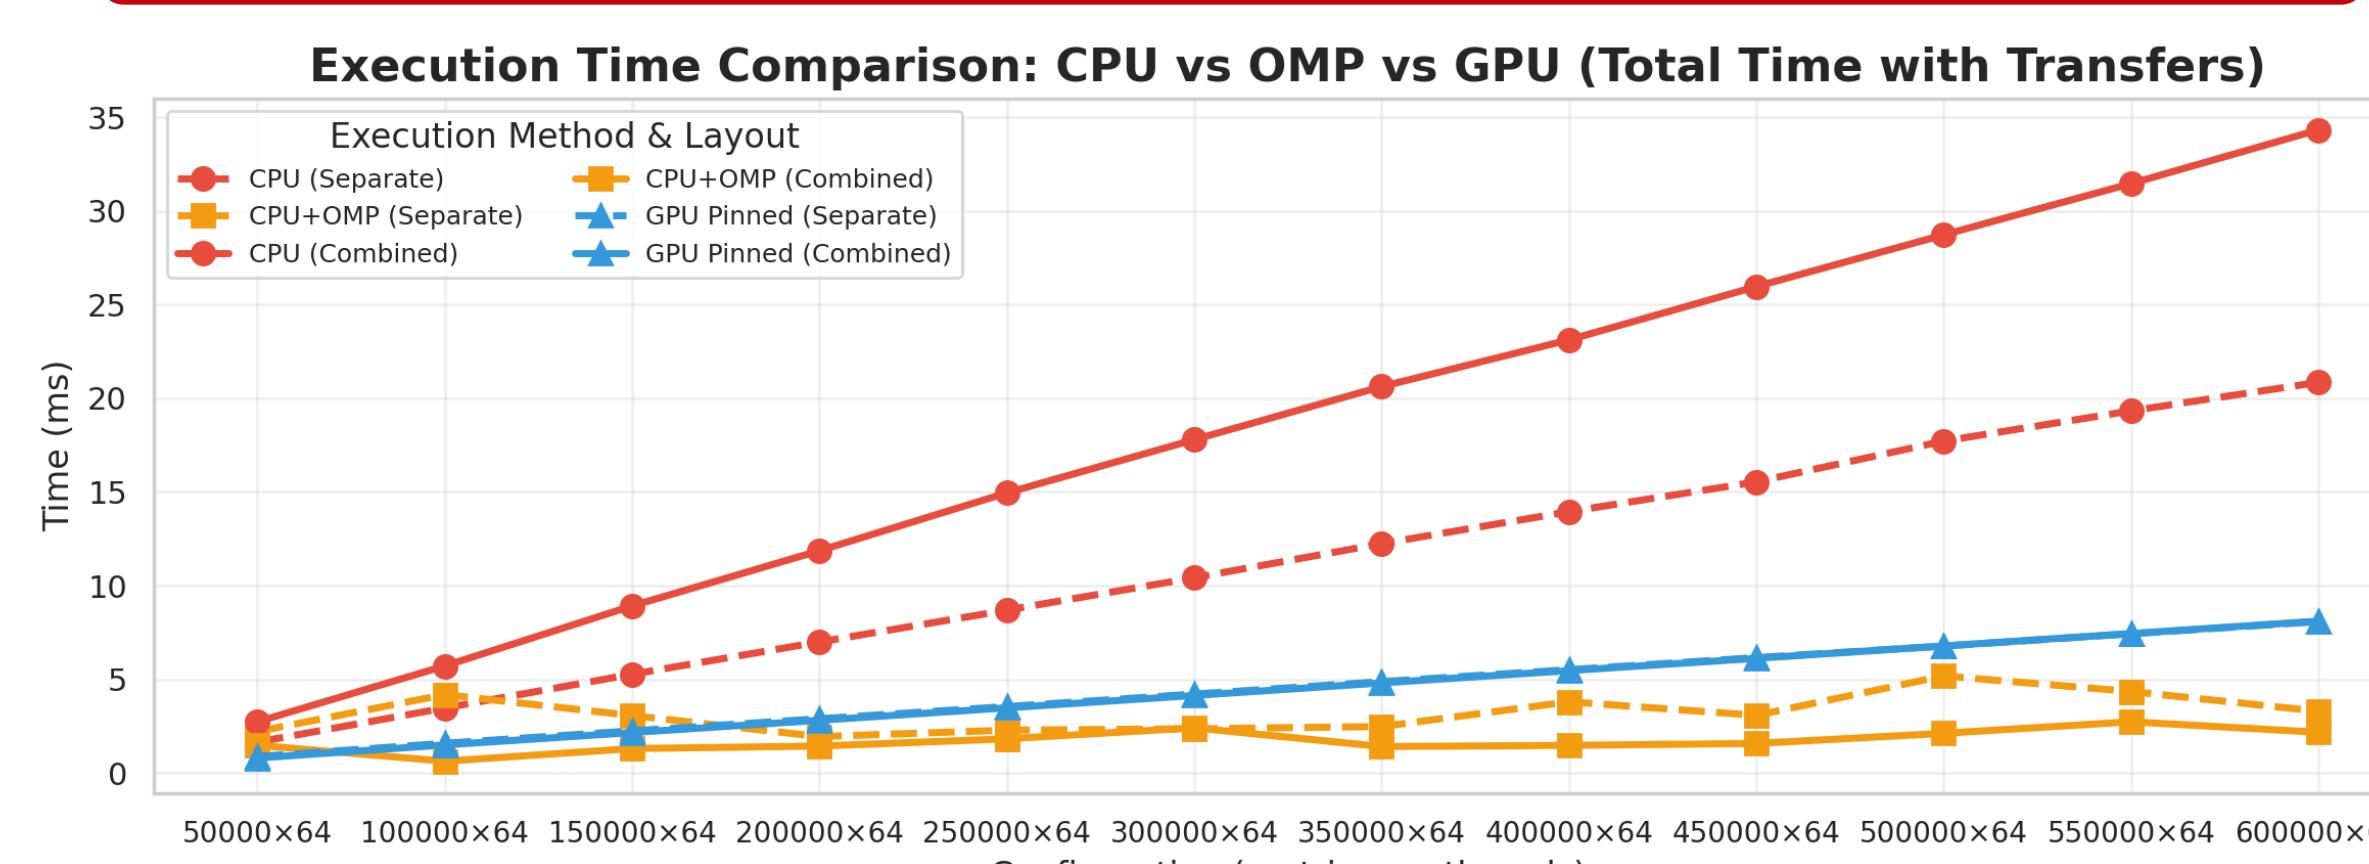


Warp Kinematics: Yucheng Huang, Wuzhen Li, Steven Li, Lynnix Zou

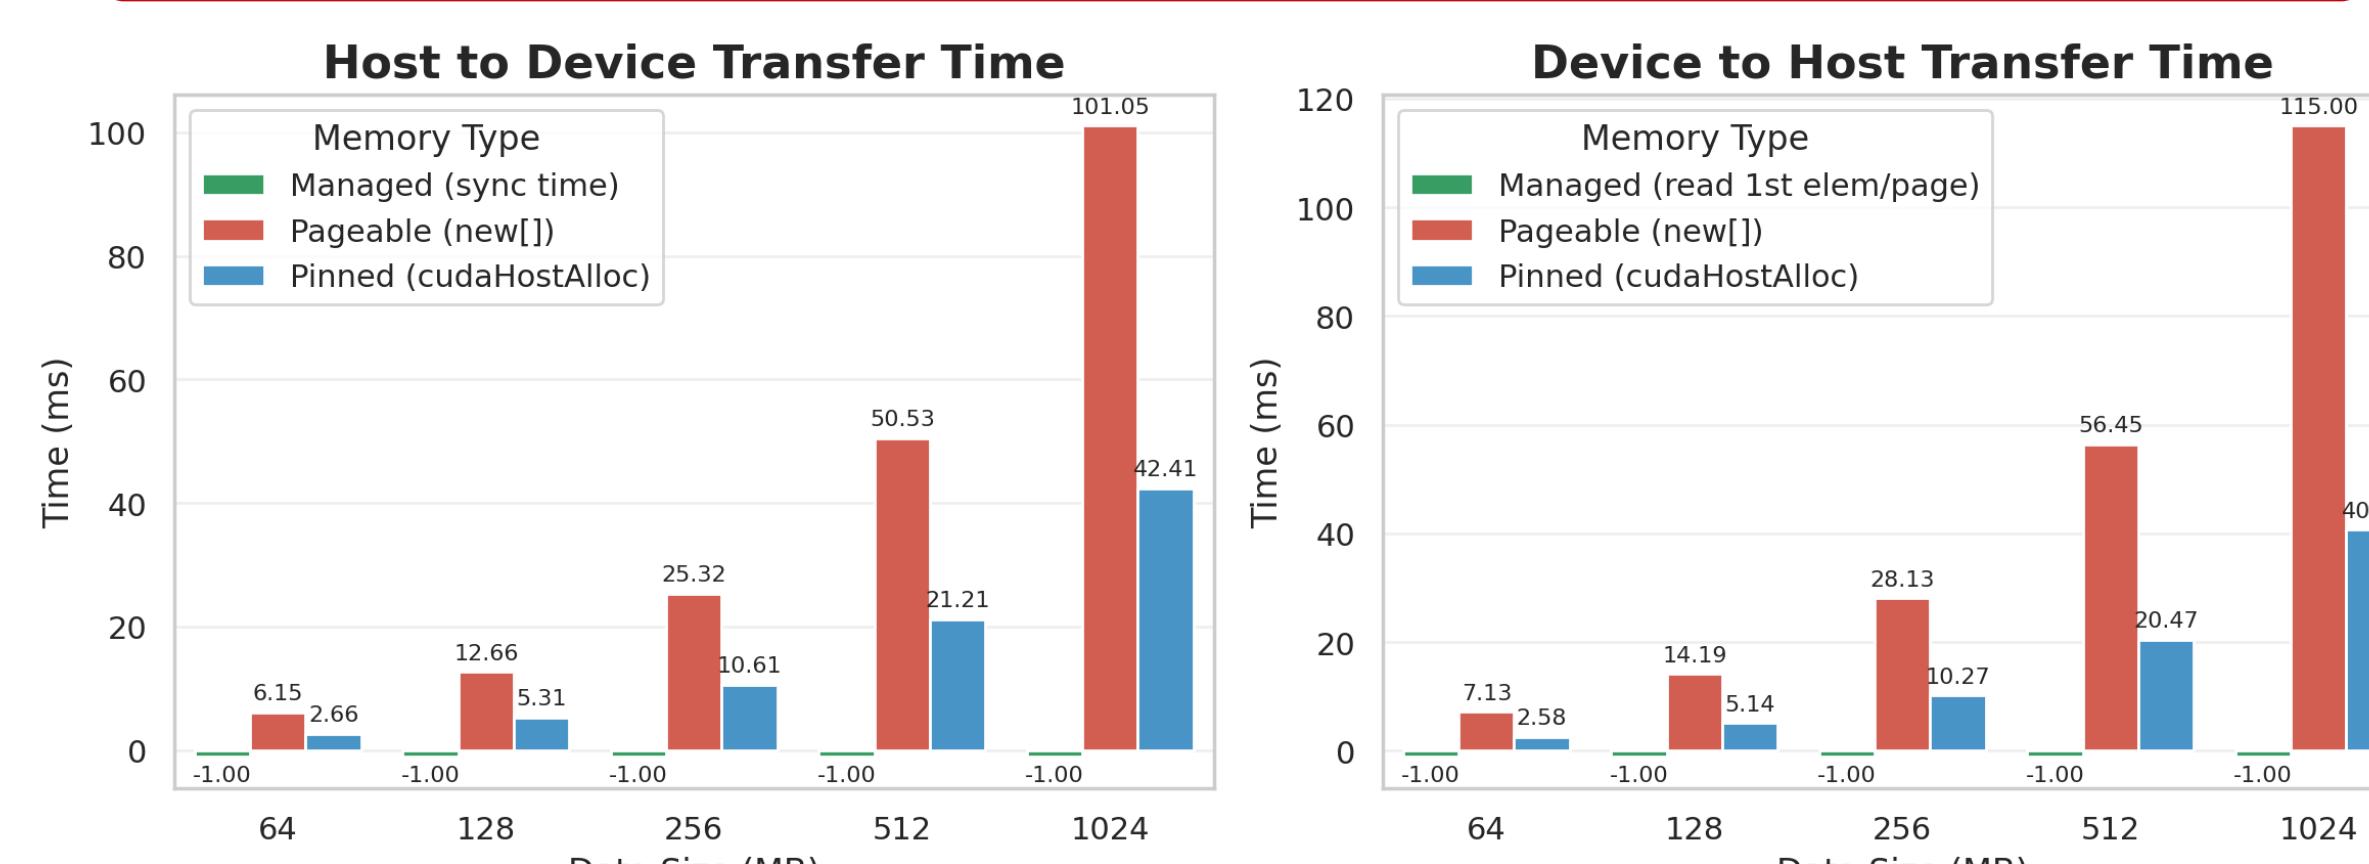
Advisor: Tsung-Wei Huang

### Workstation<sup>1</sup> - Xeon Gold 6330 + 3090

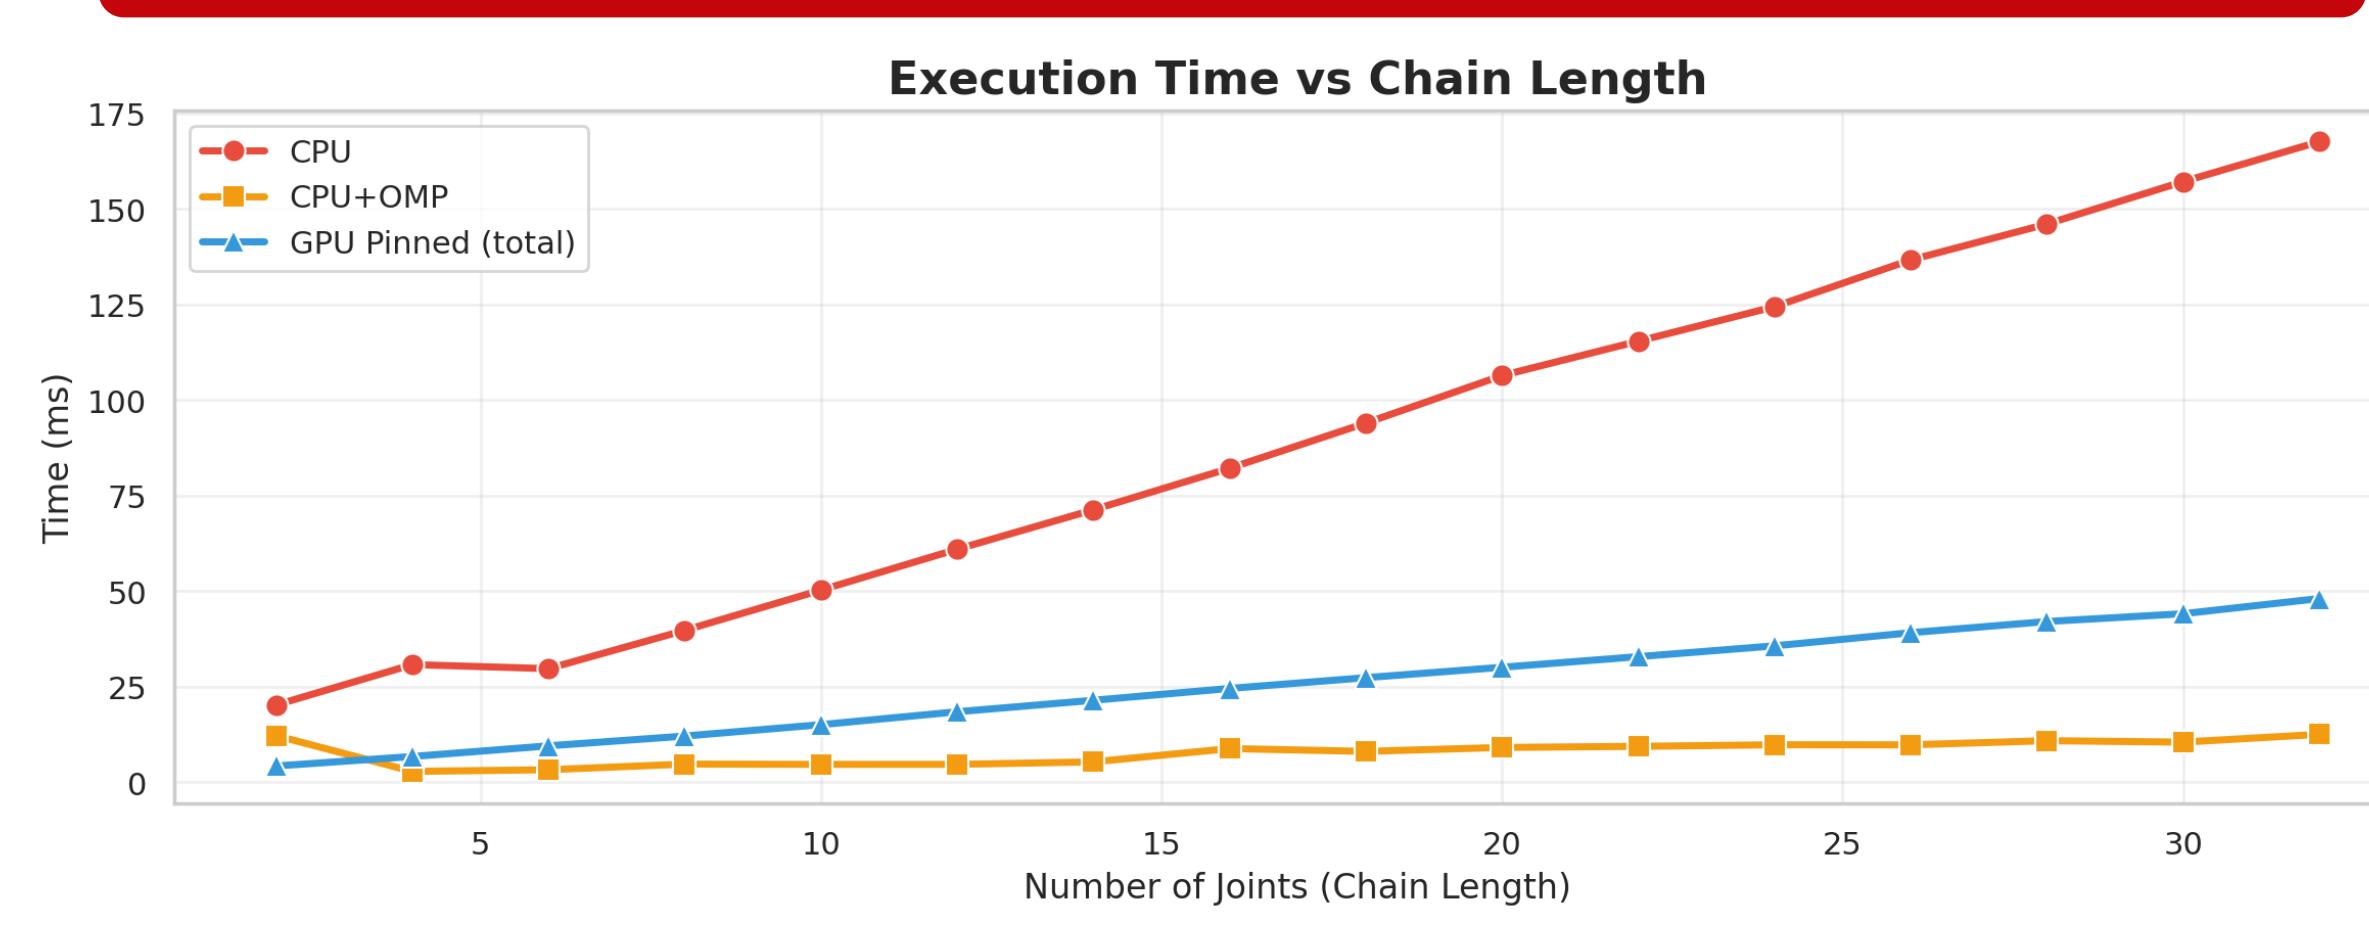
#### Data Structure Improvement



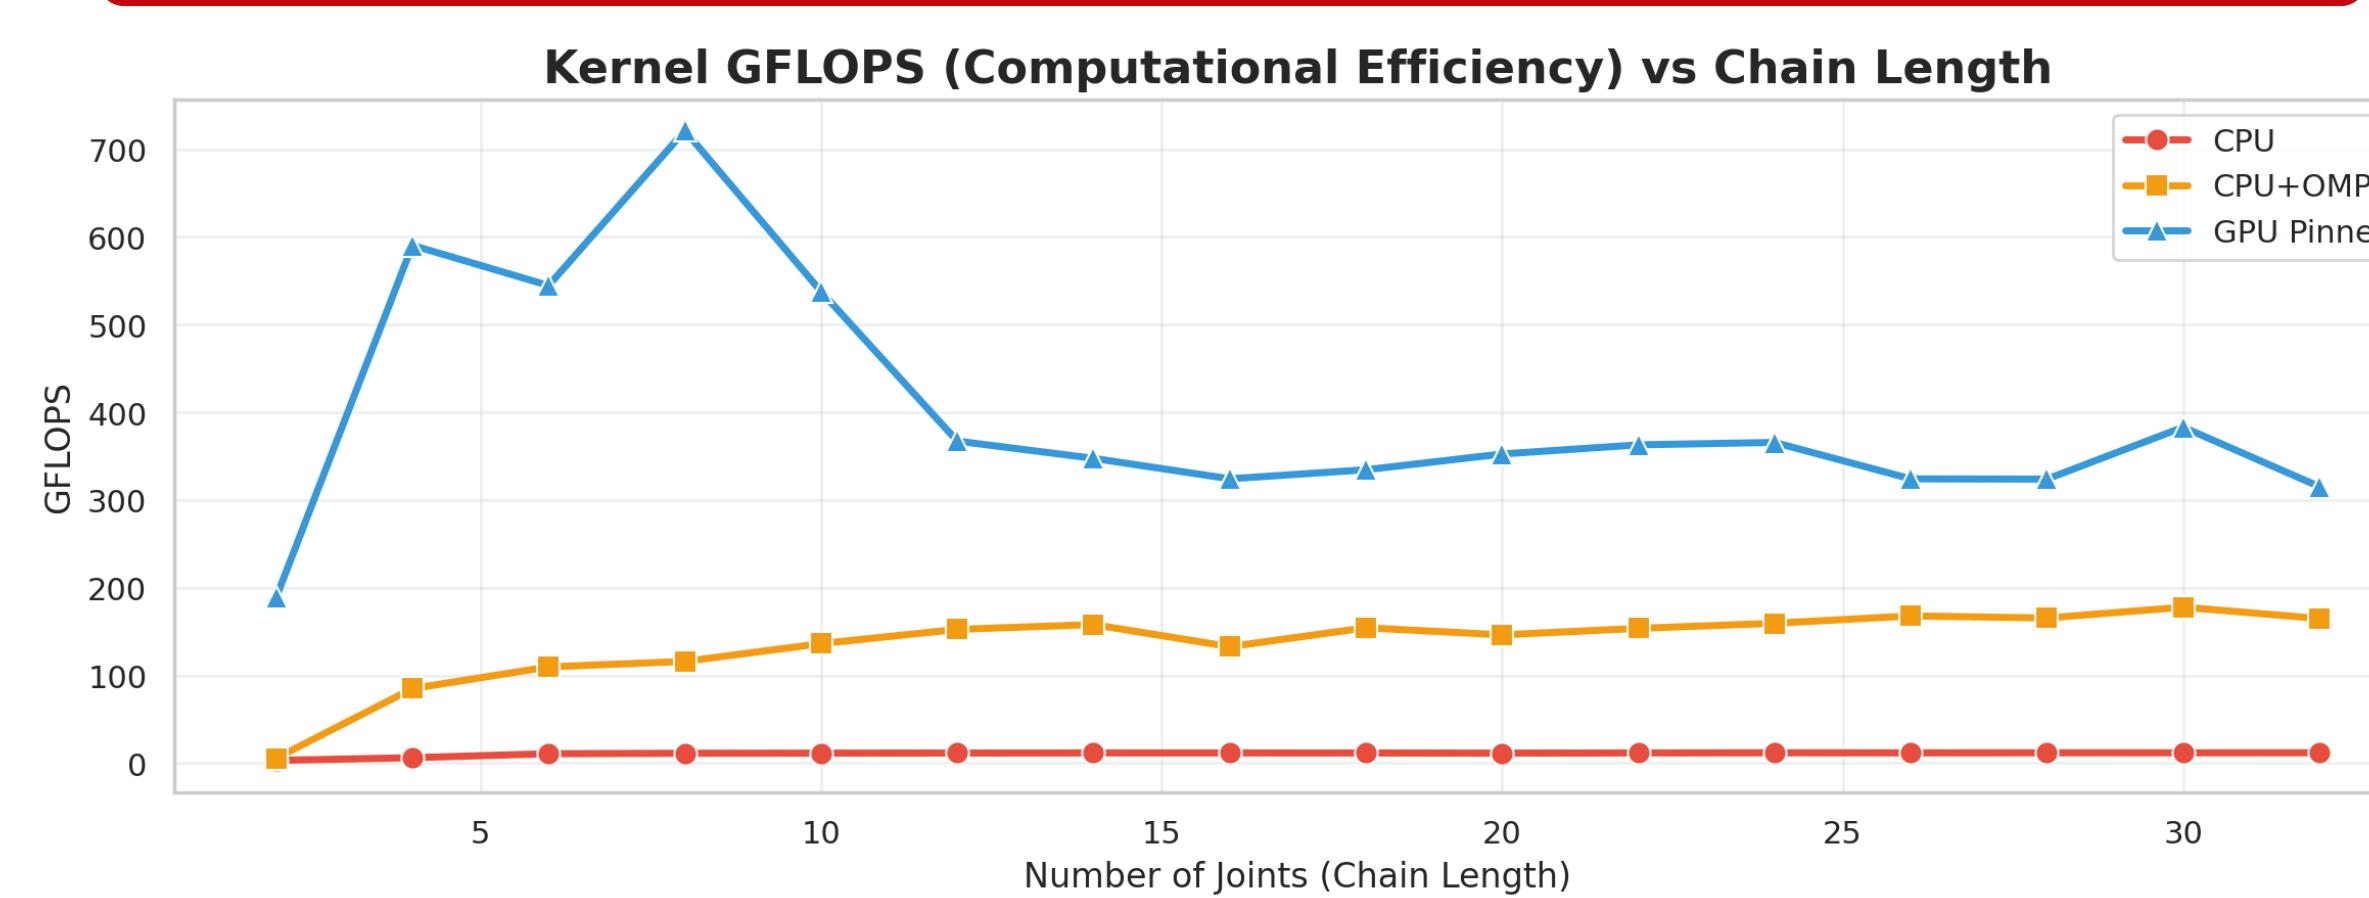
#### Data Transfer Improvement<sup>4</sup>



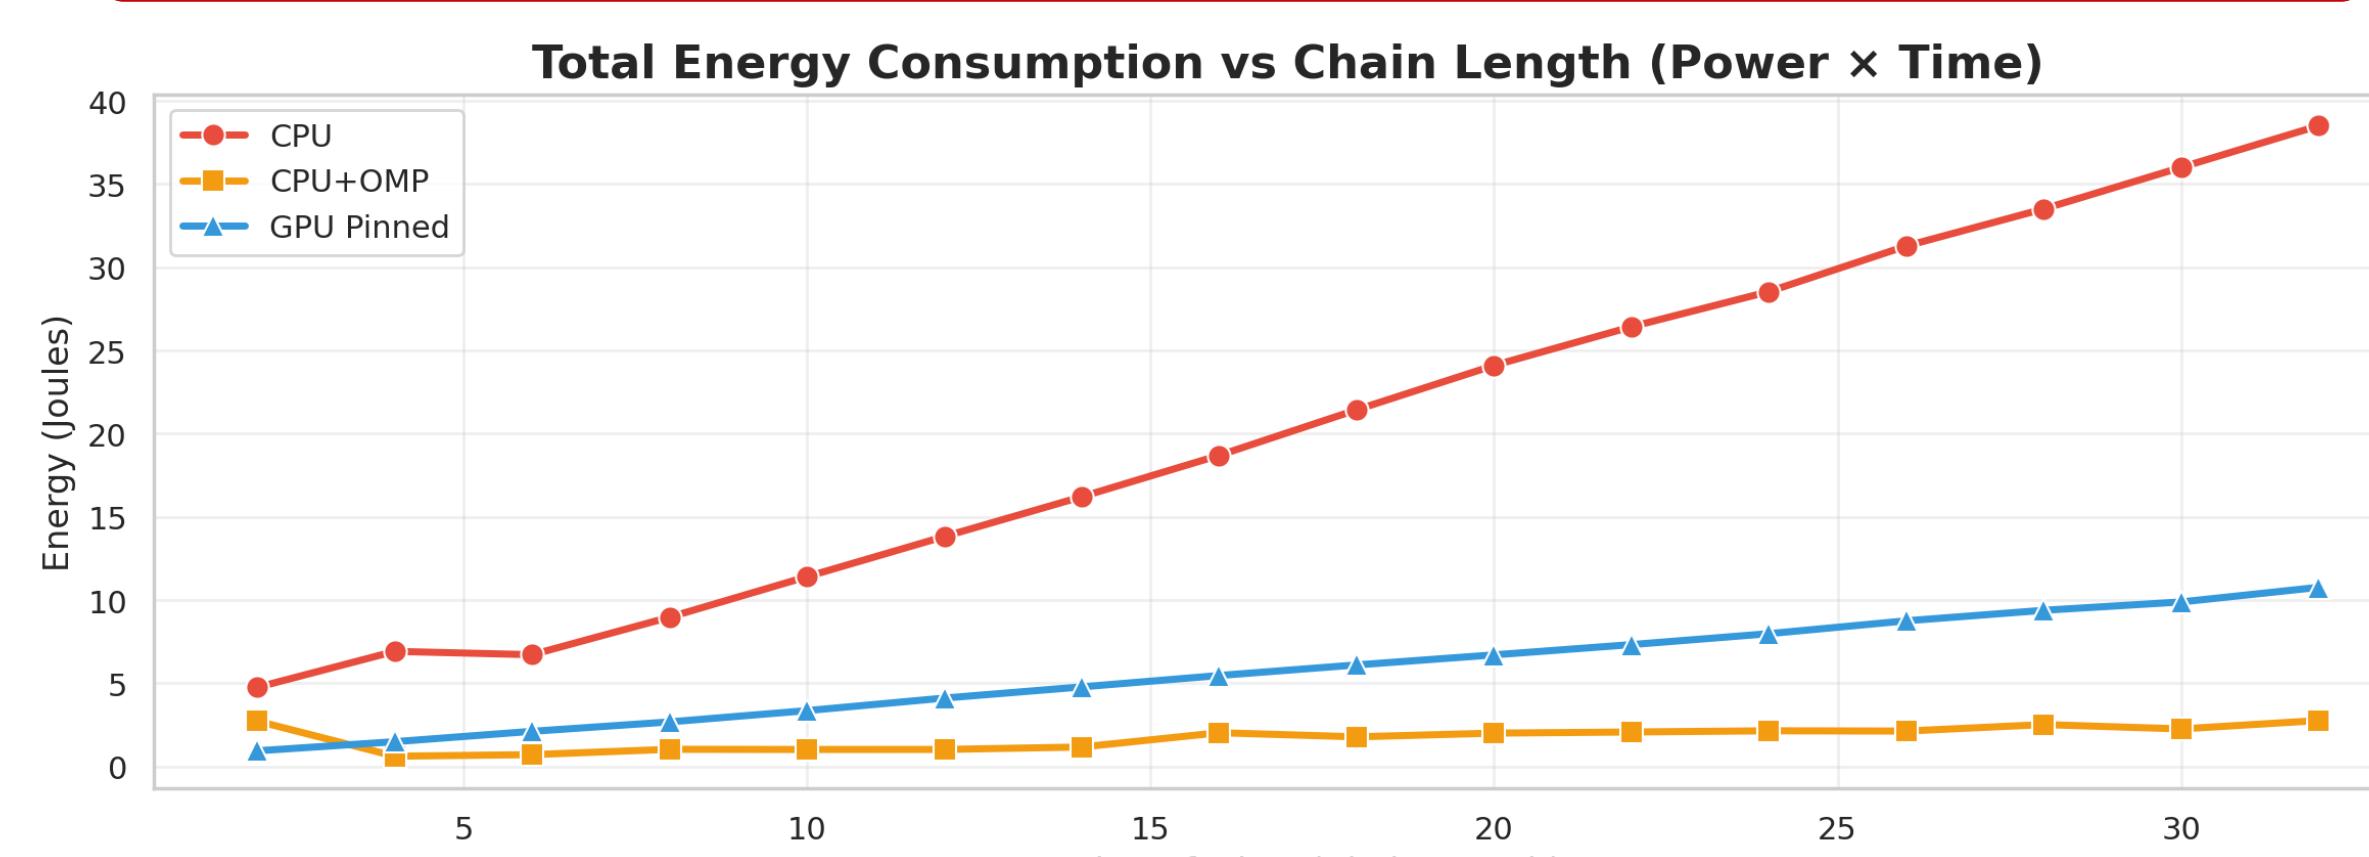
#### Variable Workload Thread Performance



#### Throughput Performance



#### CPU + GPU Energy Consumption



### Motivations and Goal

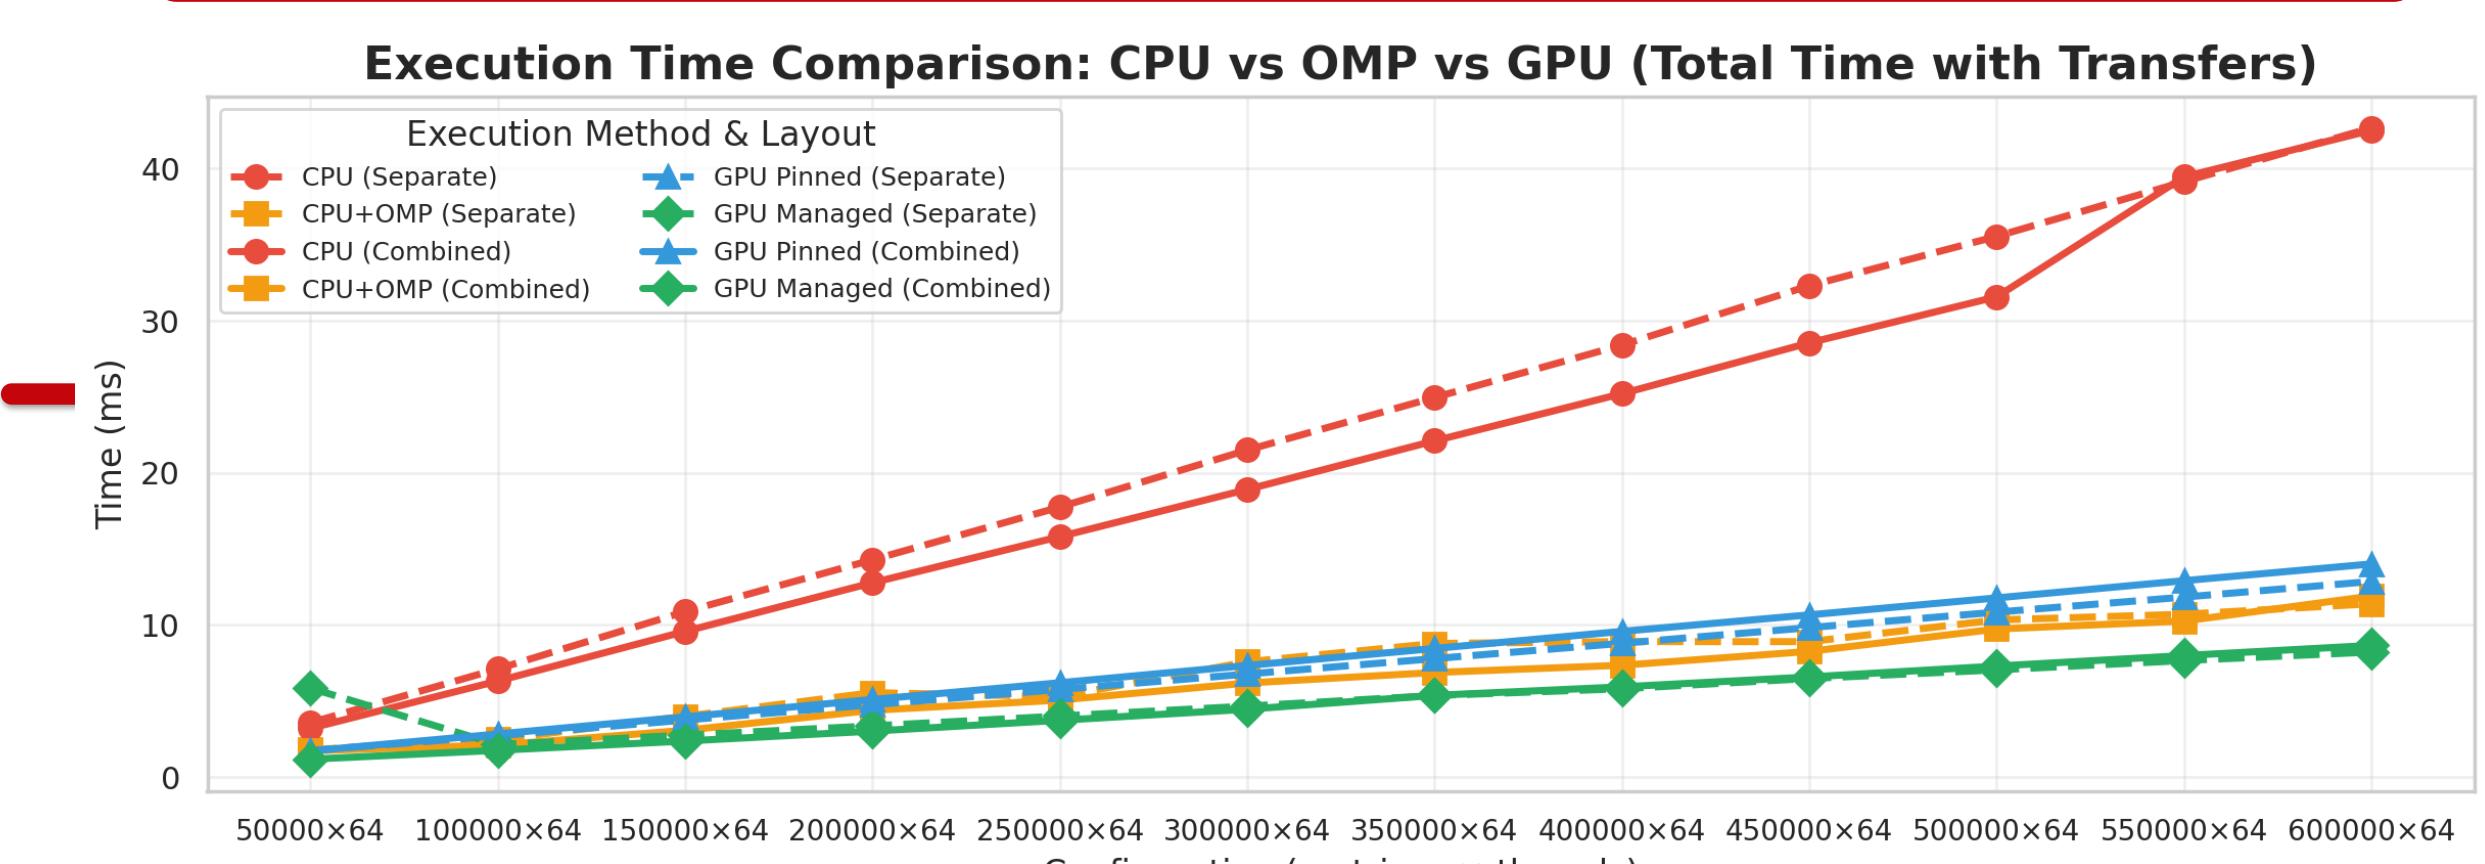
Modern robots relies on thousands of **Forward Kinematics** per move instructions using 4x4 matrix multiplications

On MCUs with limited cores, this yields poor throughput but on edge devices with CUDA we get to utilize more threads and unified memory

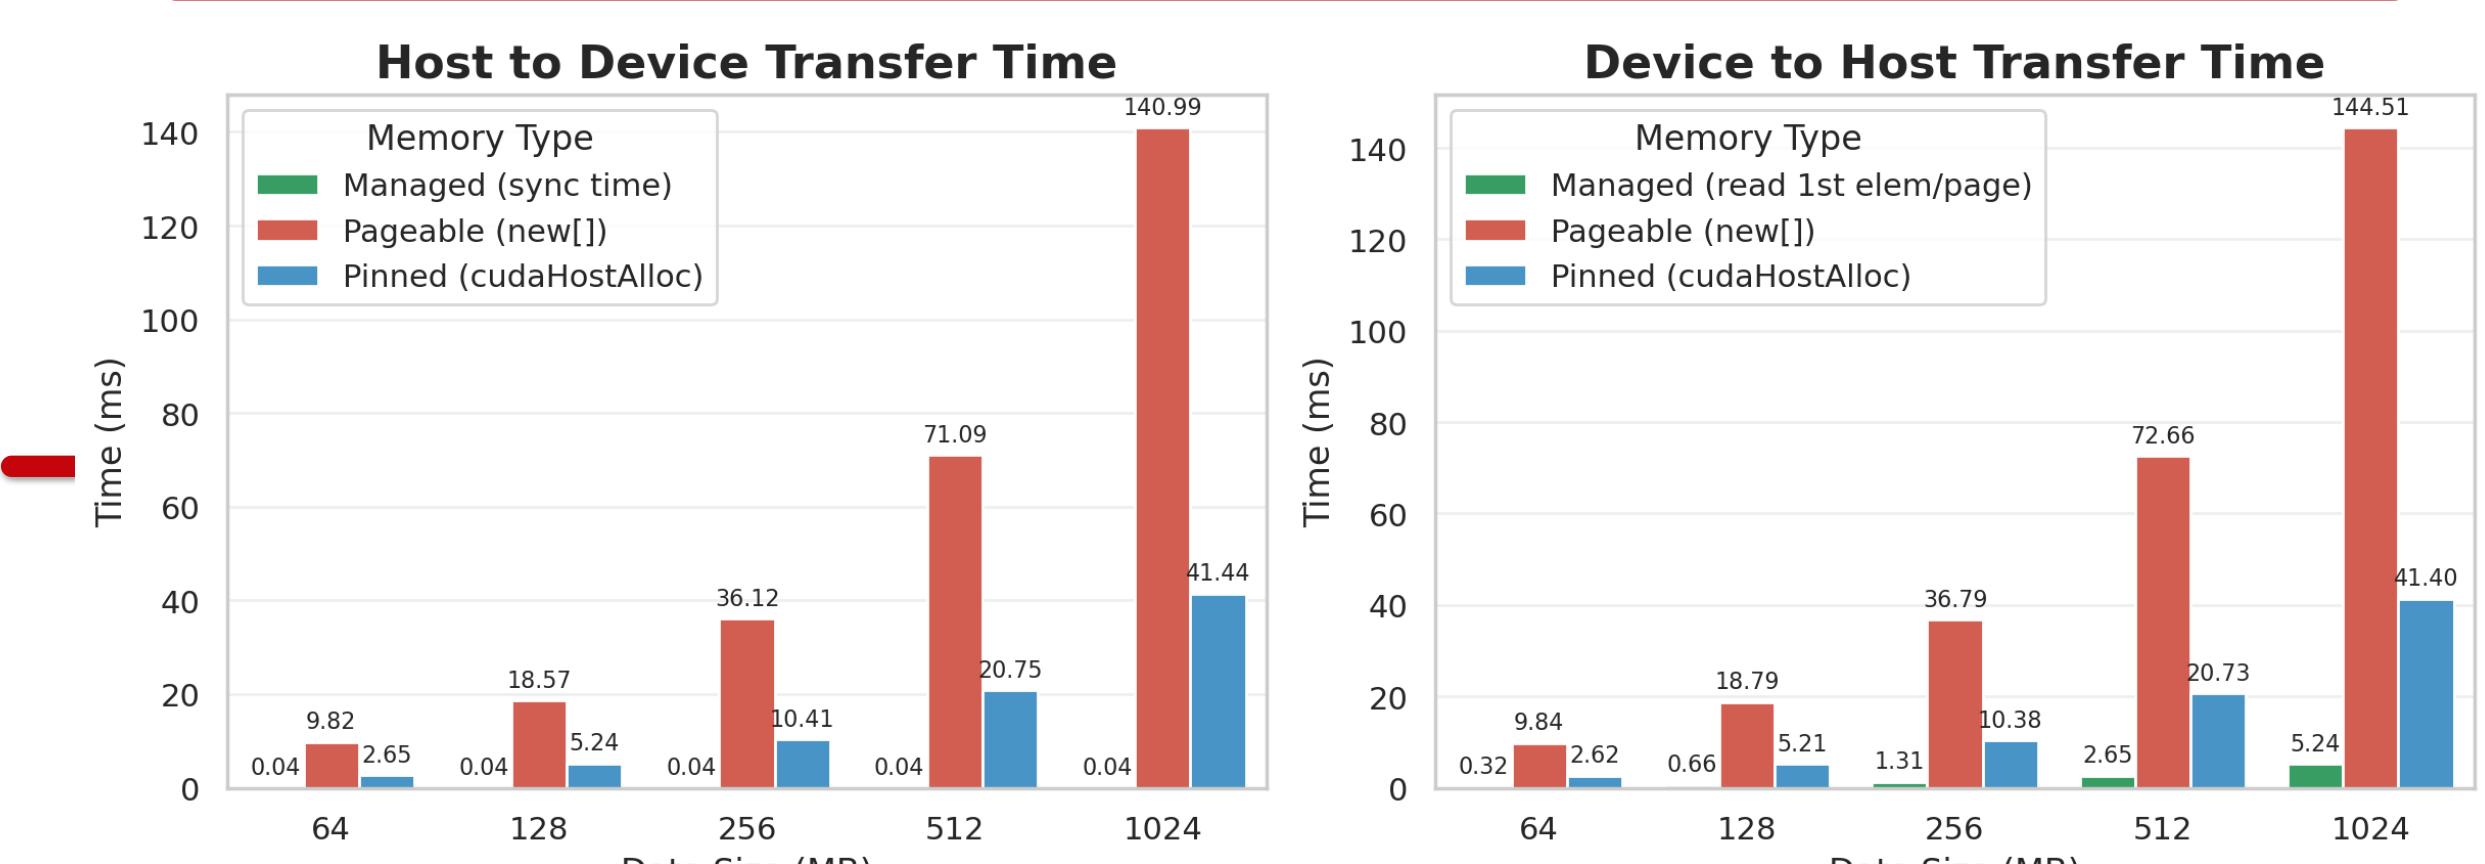
Our goal is to investigate that under what certain setups and tasks, does it worth it to compute with CUDA cores instead of CPU cores to solve Large-Scale Small-Size Matrix Multiplication problem

### Edge Device<sup>2</sup> - Jetson Nano Super

#### Data Structure Improvement



#### Data Transfer Improvement



### Methodology and Experiments<sup>3</sup>

#### Data Layout Strategy: Pinned Memory Comparison:

**Method A:** Separate Memory Chunks (Per-step separation)

**Method B:** Continuous Interleaved Chunk

#### Key Finding:

**Discrete Memory:** Continuous chunks cause poor prefetching behavior for Single Core CPU version

**Unified Memory:** Separate and continuous memory chunks perform roughly the same due to conservative prefetching

#### Platform

Edge Device

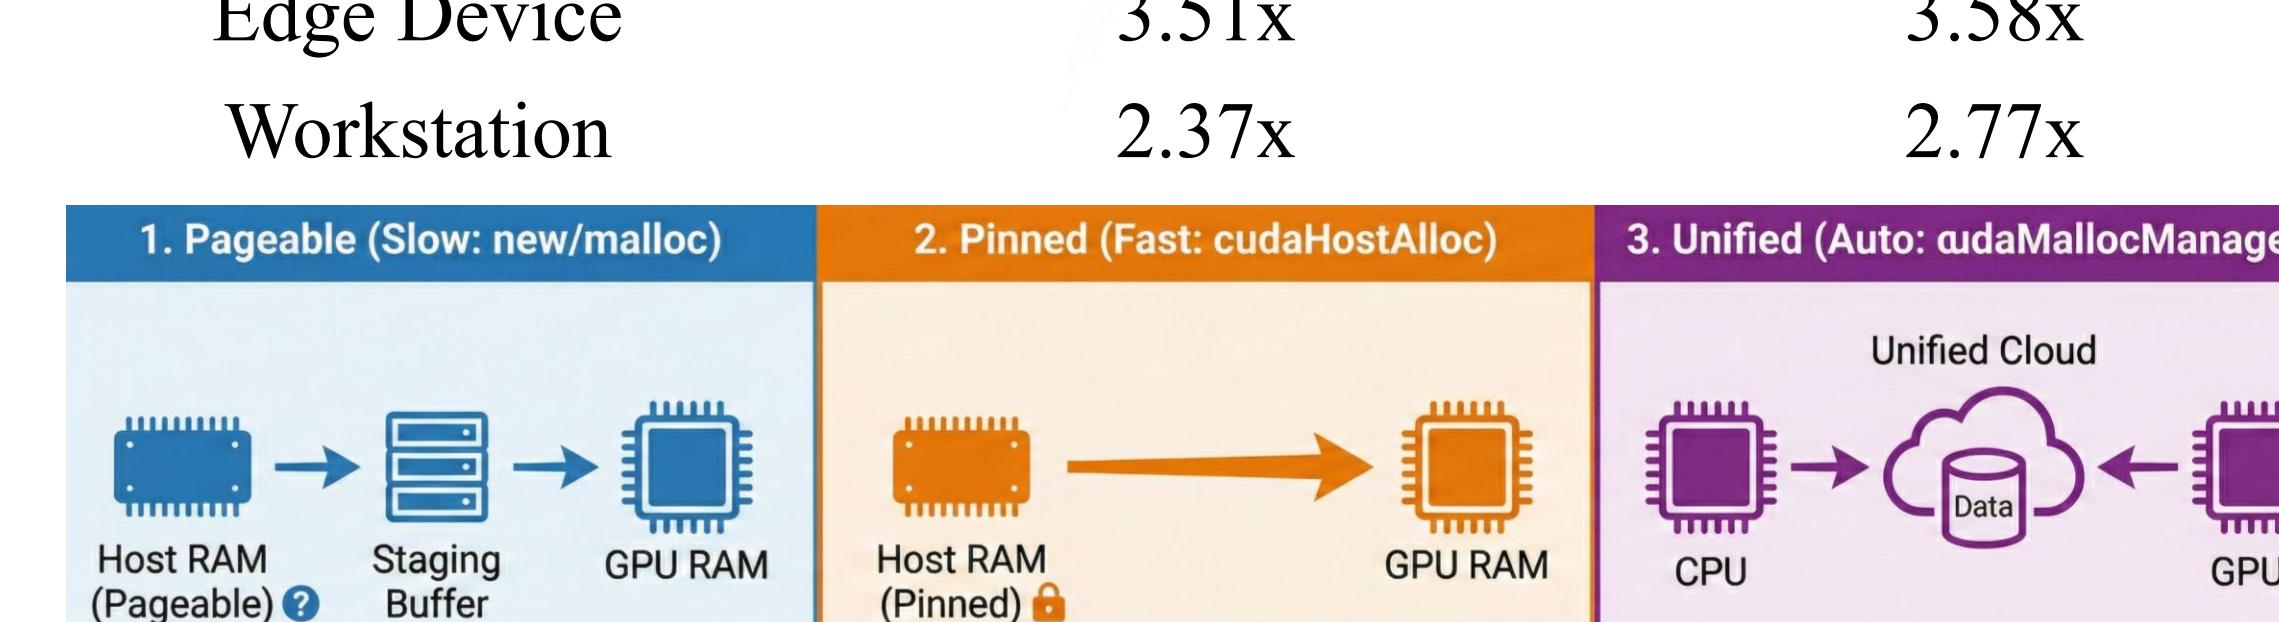
#### Host to Device

#### Device to Host

Workstation

#### Host to Device

#### Device to Host



#### Discrete Memory:

CPU+OMP is the best solution because it has moderate level of compute power while free from PCIe data movement constraint

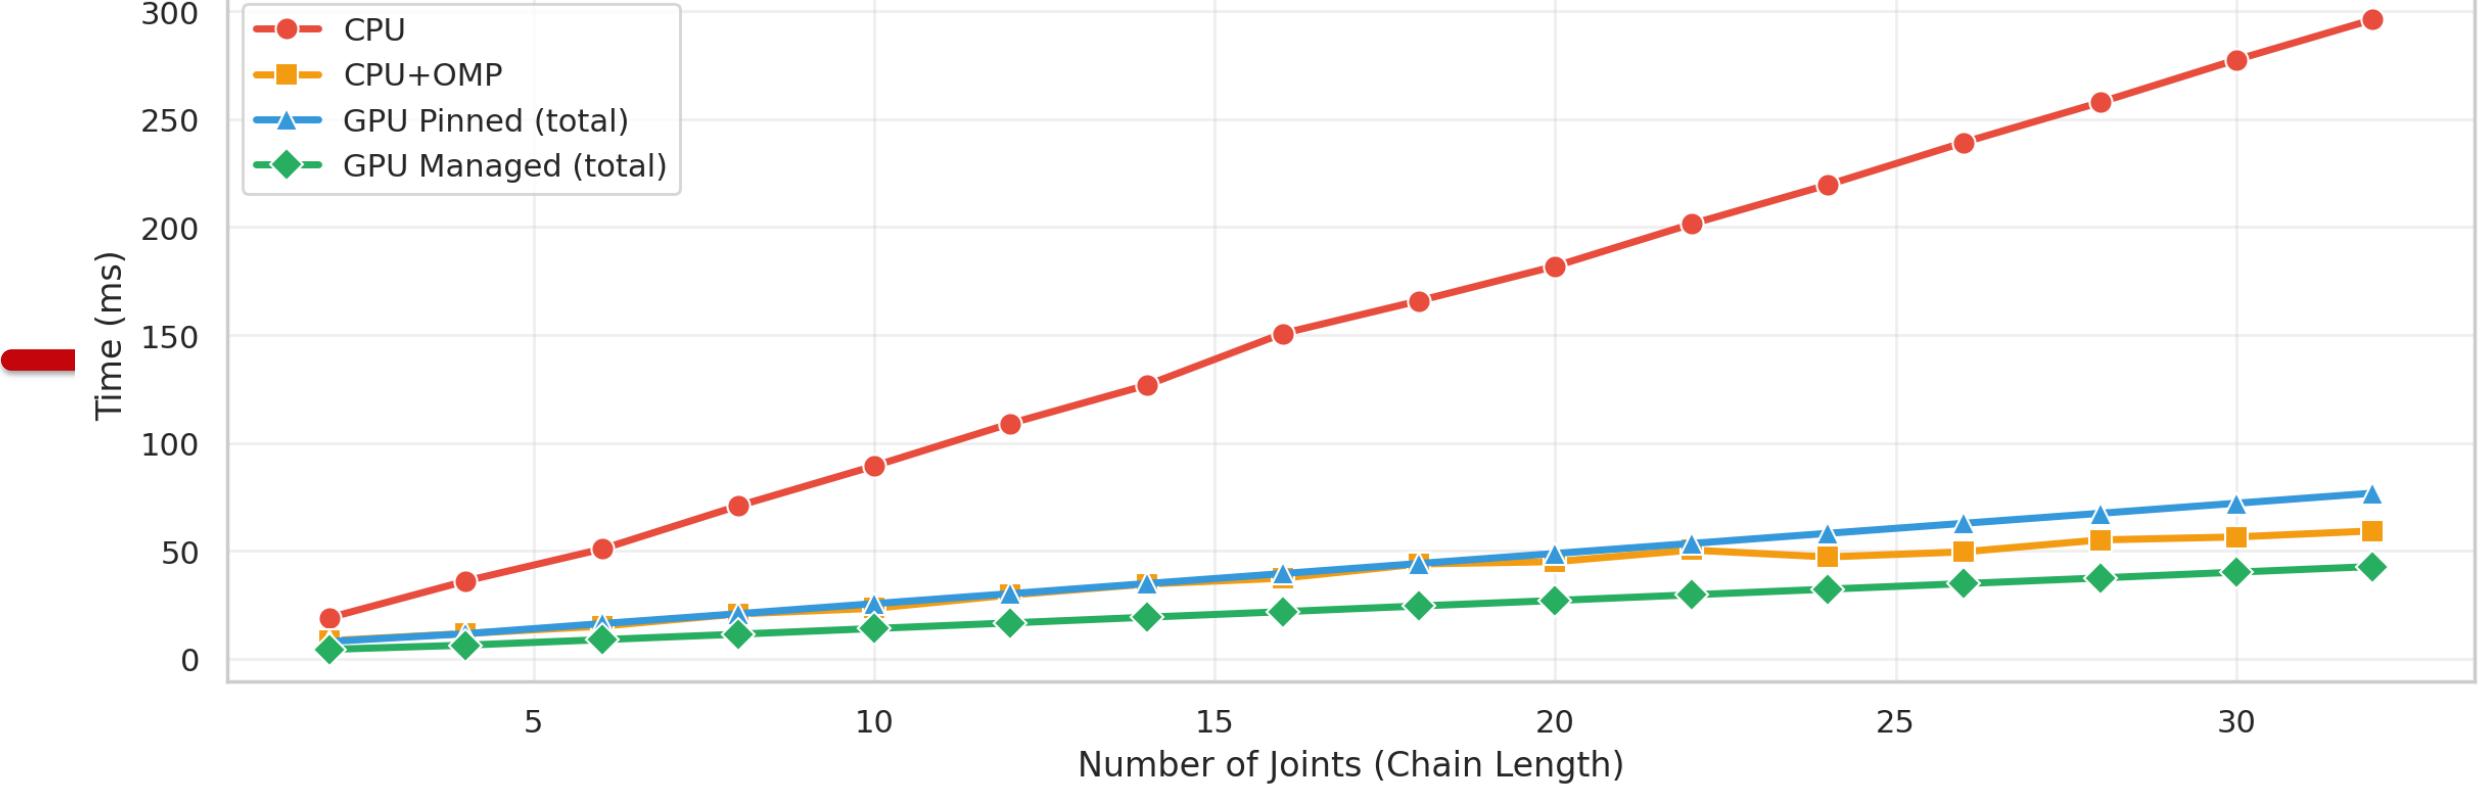
#### Unified Memory:

GPU with managed memory allocation is the best solution because it enables zero-copy compute between GPU and CPU

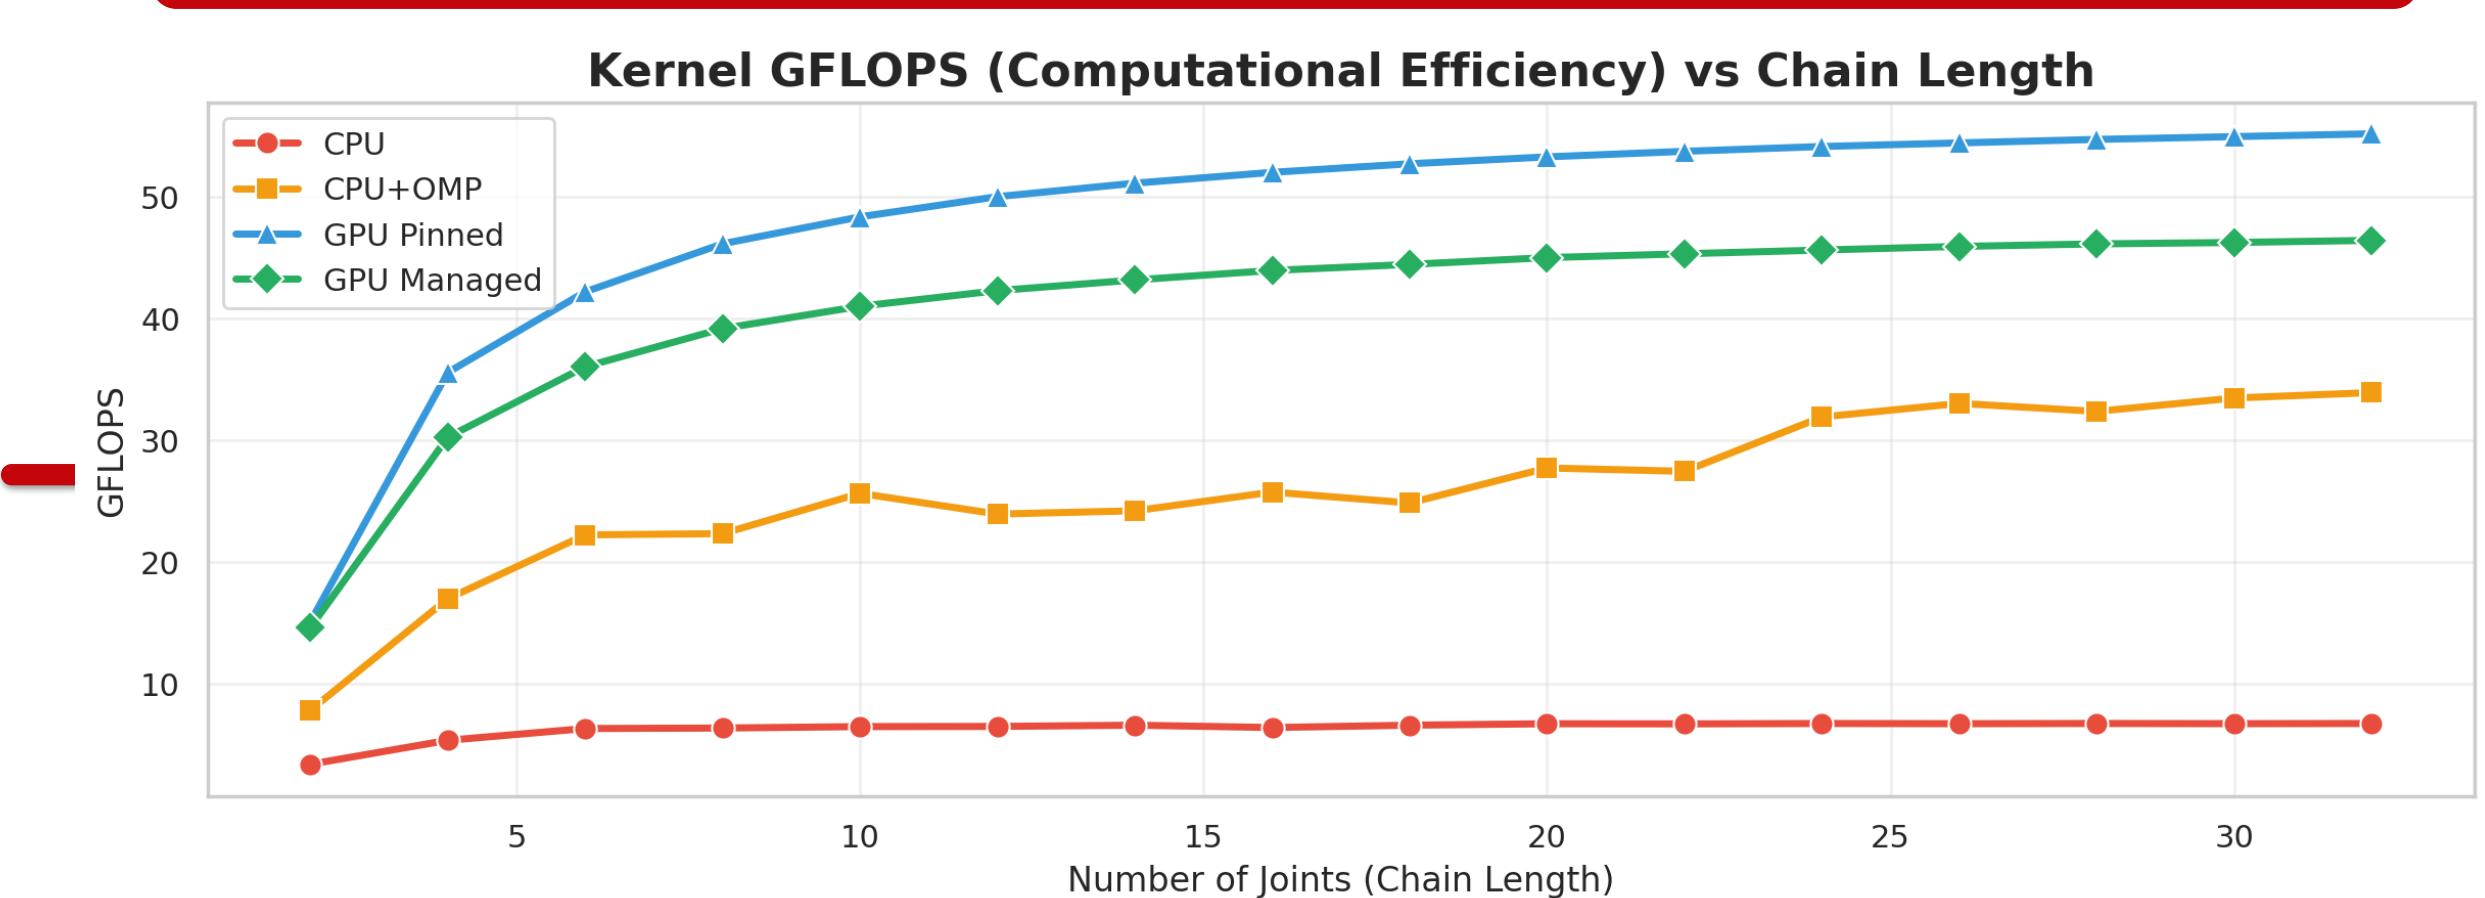
#### Device Energy Consumption (CPU OMP vs. GPU managed):

With 2.3% (0.04/1.75) the energy consumption of workstation, Jetson can process same amount of workload as the workstation at a cost of 2.96x (23.25/7.85) more time

#### Variable Workload Thread Performance



#### Throughput Performance



### Conclusion and Limitations

- Large-scale small-size matrix multiplication is best using
  - OMP on a Discrete Memory Device**
  - CUDA with pinned memory on a Discrete Memory Device with fast data transfer**
  - CUDA with managed memory on a Unified Memory Device**
- This optimized kernel does not benefit from using **CUDA Stream** because invoking a stream has additional overhead but our execution time for the kernel is too short to benefit from overlapping compute and transfer

#### CPU + GPU Energy Consumption

